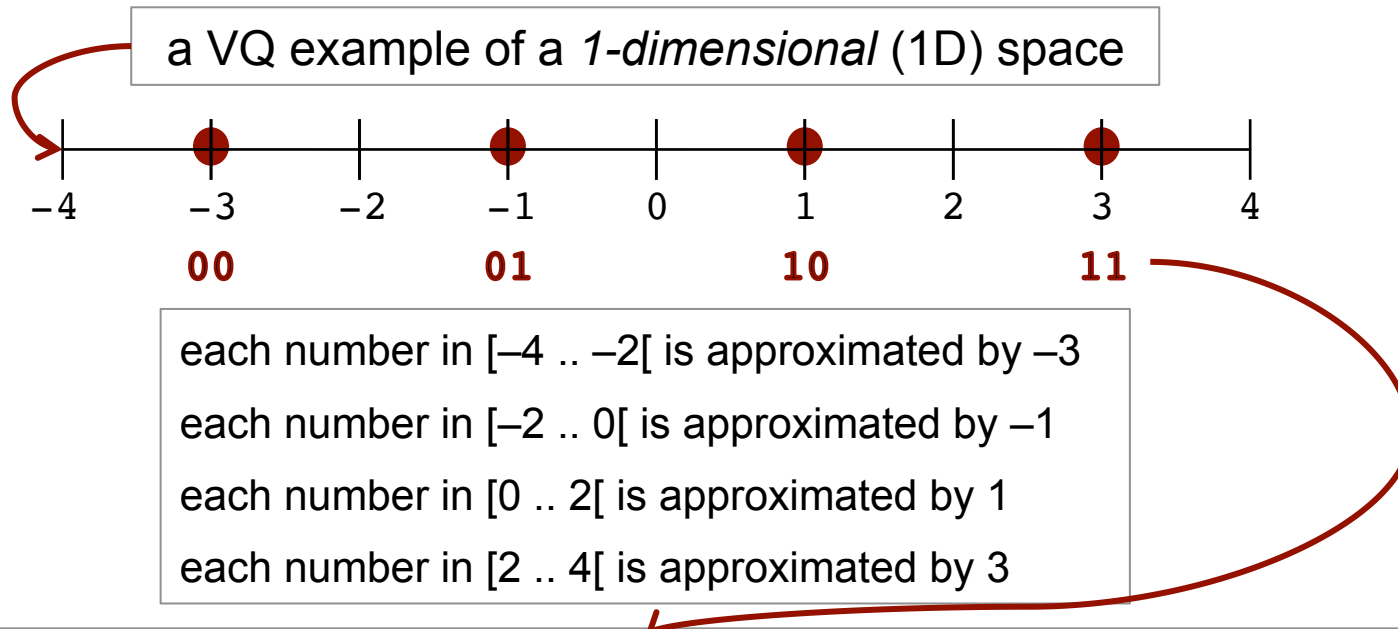


# Learning Vector Quantization (LVQ) – Neural Network

## About the “vector quantization” concept

a “vector quantization” (VQ) is nothing more than an approximator; the idea is similar to that of “rounding-off” (e.g., to the nearest integer); the original motivation is *dimensionality reduction* or *data compression*




the approximate values are uniquely represented by 2 bits

00 for  $-3$ , 01 for  $-1$ , 10 for  $1$ , 11 for  $3$

so, the approximation is also a “compressed” way of representing the data; we use 2 bits instead of the 4 bits that would be needed to represent the 9 different integers

... the “vector quantization” concept – a 2D example

a VQ example of a 2-dimensional (2D) space



each pair of numbers (i.e. each 2D point) falling in a particular region is approximated by the red cross associated with that region;  
notice that there are 16 regions and 16 red crosses each of which can be uniquely represented by 4 bits; thus, this is a 2-dimensional, 4-bit VQ

## ... “vector quantization” (VQ) and Learning VQ (LVQ)


so, the general formulation is that “vector quantization” maps *k-dimensional* vectors into a finite set of vectors  $Y = \{y_i: i=1, 2, \dots, n\}$ ;

each vector  $y_i$  is called a “codebook vector” (cBv), and the set,  $Y$ , of all cBv is called a “codebook”

the “Learning Vector Quantization” (LVQ) goal is that:

given a dataset (training sequence) and given a number of “codebook vectors”, find a codebook which best represents the class values in the dataset

*example:* each + is a training example (from the dataset) and each red-circle is a cBv  
the lines represent the decision boundaries




# LVQ representation – Neural Network (no Hidden Layers)

*one layer of:  
INPUT  
nodes*

*one layer of:  
OUTPUT  
nodes*

*one INPUT node for:  
each feature (i.e., dataset column)*



*the OUTPUT layer is also called;  
a “codebook”*

*each OUTPUT node is also called:  
a “codebook vector”*

*at least one OUTPUT node for:  
each class value*

*one weight (connection) from  
each INPUT node to each  
OUTPUT node*

... there are as many **input** nodes as  
**features** in the dataset


... there is at least one **output** node for  
each **class** value

here, D input and M output nodes; so D  
features and if using 2 nodes for each  
class value we have M/2 class values

## Example – dataset and “codebook vectors”

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

dataset



features: X1, X2  
class: Y  
class values: 0, 1

Codebook Vectors		
X1	X2	Y
3,58229404	0,79163723	0
7,79278348	2,33127338	0
7,93982082	2,86699026	1
3,39353321	4,67917911	1

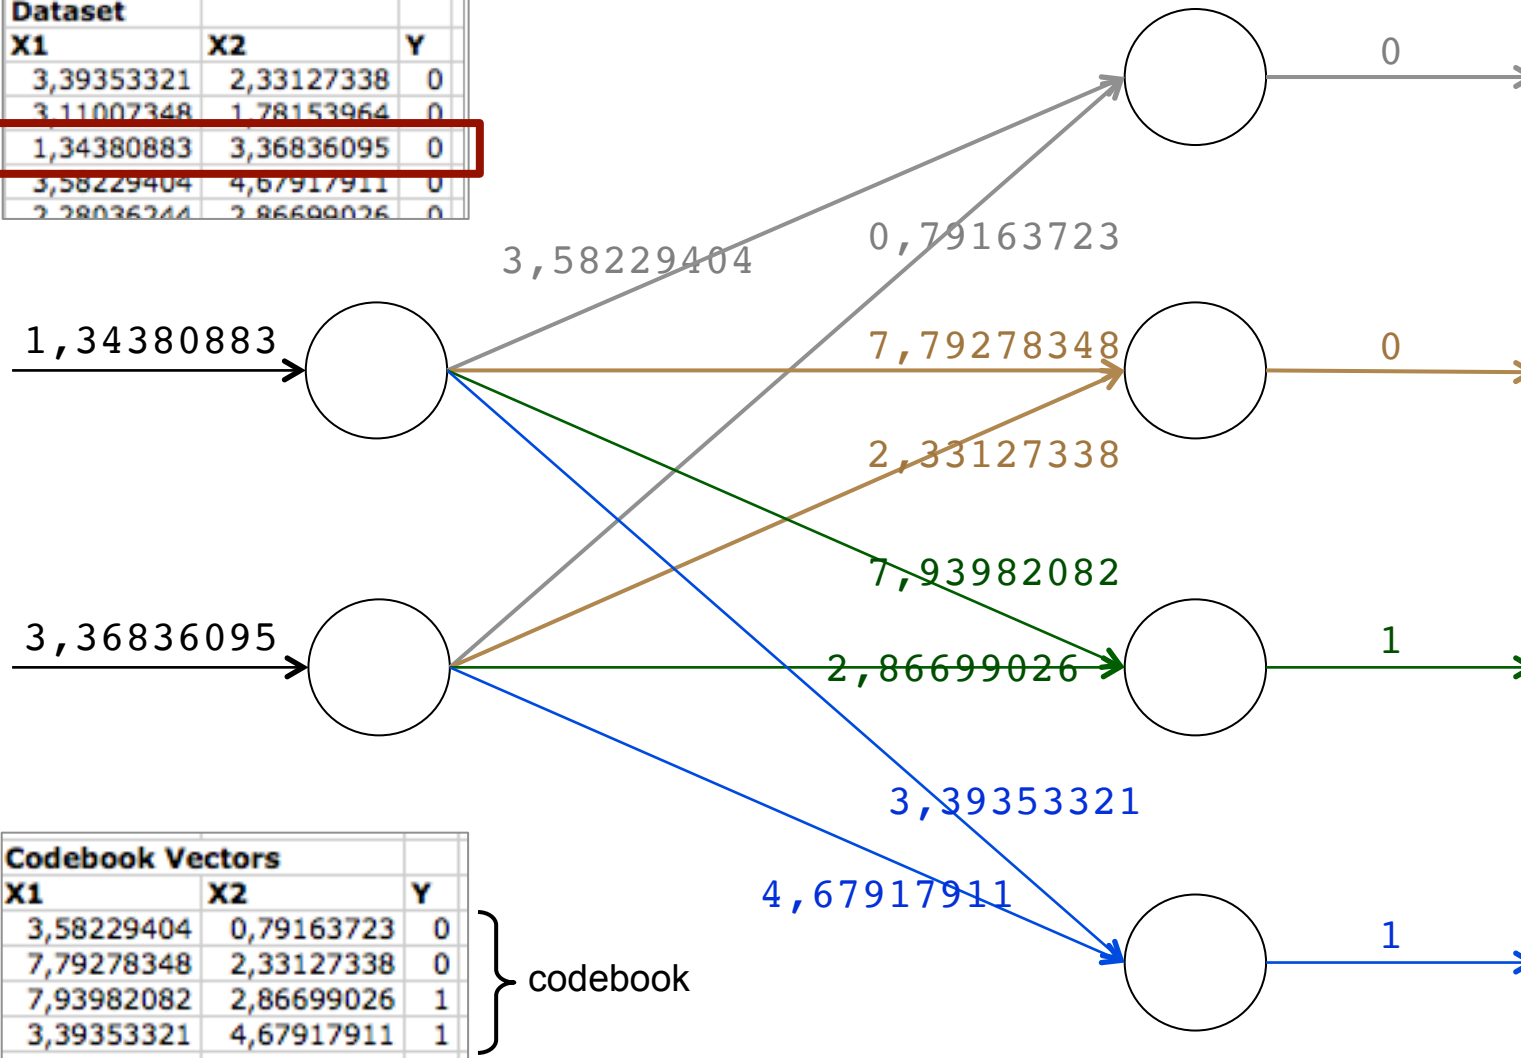
### Exercise:

given these dataset and codebook vectors draw the LVQ neural network representation;


tag the connections using the weights taken from these codebook vectors and, as input, consider the third (3<sup>rd</sup>) instance of the dataset.

## ... example – the LVQ Neural Network representation

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0



... example – plot of dataset and the initial codebook



Codebook Vectors		
X1	X2	Y
3,58229404	0,79163723	0
7,79278348	2,33127338	0
7,93982082	2,86699026	1
3,39353321	4,67917911	1

} the **initial** codebook

we defined 2 codebook vectors (cBv) for each class value;

i.e, 2 cBv for class value 0 and 2 cBv for class value 1



## LVQ method – a “competitive learning” approach

*... the intuition is that,*  
given a **training example**, each output unit (codebook vector)  
“competes” with all other output units

*and, the winner is,*  
the output unit **closer** to that training example

*and, the winner gets (as a reward):*  
an **update of its weights** aiming to get more adapted than other units;  
the adaptation is processed via **attraction** and **repulsion** rules

**attraction** rule – move codebook vector **closer to** the training example  
**repulsion** rule – move codebook vector **away from** the training example

*because one codebook vector is selected for modification for each training instance,  
the algorithm is referred to as a winner-take-all (type of competitive learning)*

## The LVQ learning process – from the Data to the Model

the LVQ algorithm learns the “codebook vectors” from the training data

- **[1]** choose the number of codebook vectors (cBv) to use
  - possibly (but not necessarily) an equal number of cBv for each class
- **[2]** start the learning process with a pool of cBv
  - either randomly selected from training data,
  - or randomly generated with the same scale as the training data
  - ... each cBv has the same number of attributes as data and an output class value
- **[3]** the instances in the training data are processed one-at-a-time
  - for a given training instance, select its **most similar** cBv
  - ... the selected cBv is the “winner”, also called the “best matching unit” (**BMU**)
- **[4]** the BMU gets its weights updated
  - if the BMU has the same class as the training instance, apply the **attraction** rule
  - otherwise, apply the **repulsion** rule (in relation to that training instance)

“similar to” – how to compute this concept?

... for a given instance, select its **most similar** cBv

a common way of computing similarity is to establish a **distance** criteria  
and assume such **distance as the similarity order** relation  
i.e., the *closer the higher similarity*

so, we need to define&compute a **distance** between vectors  
and **rank** the distance from an instance data to each cBv  
the **closer** they are the **more similar** they are!

“similar to” – define&compute a distance criteria

*a general formulation of several distance functions  
(or, distance criteria) commonly used for numeric data*




$$d_k(\vec{x}, \vec{y}) = \left( \sum_{i=1}^n |x_i - y_i|^k \right)^{\frac{1}{k}}$$

Well-known special cases from this family are:

$k = 1$  : Manhattan or city block distance,

$k = 2$  : Euclidean distance,

$k \rightarrow \infty$  : maximum distance, i.e.  $d_\infty(\vec{x}, \vec{y}) = \max_{i=1}^n |x_i - y_i|$ .



all points lying on a circle or rectangle are sharing the **same distance to the center point** according to the corresponding distance function


“similar to” – the Manhattan distance

the Manhattan distance, or the “city block” distance  
inspired on the idea of Manhattan having a “grid format”

the Manhattan distance between

$A = \langle 0, 1 \rangle$  and  $B = \langle 3, 3 \rangle$

$$3 - 0 + 3 - 1 = 3 + 2 = 5$$



the Manhattan distance

$$ManhDist(X, Y) = \sum_{i=1}^n |x_i - y_i|$$

compute Manhattan distance between

$A = \langle 0, 3, 2, 1, 10 \rangle$  and  $B = \langle 2, 7, 1, 0, 0 \rangle$

$$|0-2| + |3-7| + |2-1| + |1-0| + |10-0| = 18$$


## “similar to” – the Euclidean distance

straight line distance between two points

the Euclidean distance between

$A = \langle 0, 1 \rangle$  and  $B = \langle 3, 3 \rangle$

$$[(3 - 0)^2 + (3 - 1)^2]^{1/2} = [9 + 4]^{1/2} = 3.6$$



the Euclidean distance between  
vectors  $d_j$  and  $d_k$

$$|d_j - d_k| = \sqrt{\sum_{i=1}^n (d_{i,j} - d_{i,k})^2}$$

compute Euclidean distance between

$d_1 = \langle a, b, c \rangle$  and  $d_2 = \langle x, y, z \rangle$

$$\sqrt{\text{Abs}[a - x]^2 + \text{Abs}[b - y]^2 + \text{Abs}[c - z]^2}$$

## “attraction/repulsion” – how to compute these concepts?

... if the BMU has *the same class* as the training instance, **then** apply **attraction** rule; **otherwise** apply **repulsion** rule

... and an additional question is:  
what is the **amount** that the vector is **moved**?  
such amount is called the “**learning rate**”


**attraction** rule – move codebook vector,  $x$ , **closer to** the training example,  $t$ , by the amount of *LearningRate*,

$$x = x + \text{LearningRate} \times (t - x)$$


**repulsion** rule – move codebook vector,  $x$ , **away from** the training example,  $t$ , by the amount of *LearningRate*,

$$x = x - \text{LearningRate} \times (t - x)$$

## “attraction/repulsion” – an illustrative example



attraction rule



repulsion rule


- $\vec{x}$ : data point,  $\vec{r}_i$ : reference vector
- $\eta = 0.4$  (learning rate)

... this would be repeated for each training instance;  
**one iteration** of the training dataset is called an **epoch**;  
the process is completed for a **number of epochs** that we define, e.g., 200



## “attraction/repulsion and epoch” – an illustrative example


each training instance originates the attraction/repulsion of only one cBv  
an **epoch** has as many cBv movements as the number,  $N$ , of training instances  
a **number of epochs**,  $\text{MaxEpoch}$ , originates  $\text{MaxEpoch} \times N$  movements of cBv




in this illustration we have one  
codebook vector (cBv)  
associated with each class value;  
but we may have several cBv  
associated with each class value

the amount that each cBv  
is moved is controlled by  
the *LearningRate*;  
in this illustration we have a  
**fixed** *LearningRate* = 0.1

## The *LearningRate* – fixed or time-dependent?



*problem* – a **fixed** learning rate  
the process may get trapped  
into an oscillatory mode



*solution* – a **time-dependent**  
learning rate  
the rate decreases with **epoch**

## The *LearningRate* – a time-dependent approach

$$LearningRate = alpha \times (1 - \frac{Epoch}{MaxEpoch})$$

*LearningRate* is the learning rate for the current *Epoch* (from 0 to *MaxEpoch* – 1)

*alpha* is the learning rate specified to the algorithm at the start of training run

so, as *Epoch* increases (i.e., as time goes) the *LearningRate* decreases, and therefore the attraction/repulsion movements get smaller and smaller

using the above approach the *LearningRate* decreases linearly with the *Epoch*

but, there are other alternatives to define a time-dependent *LearningRate*, such as,

$$LearningRate = alpha^{Epoch}, \text{ with } 0 \leq alpha \leq 1$$

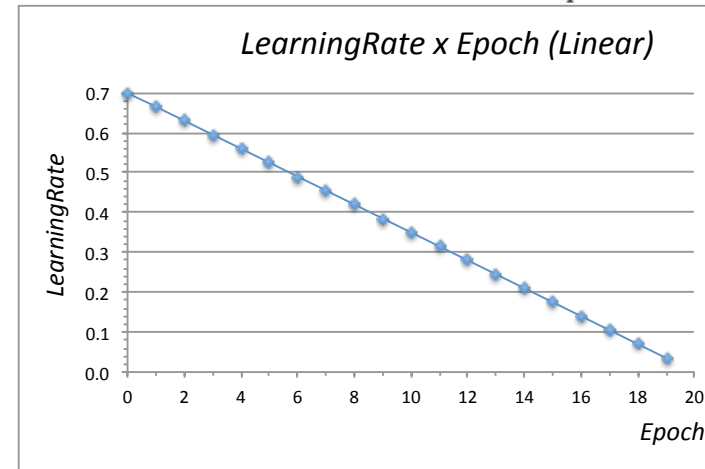
where the *LearningRate* decreases in a non-linearly with the *Epoch*

# The time-dependent *LearningRate* – an example


$\alpha = 0,7$   
 $MaxEpoch = 20$

	<i>Epoch</i>	<i>Linear LearningRate</i>	<i>non-Linear LearningRate</i>
$\alpha$	0	0,70000	1,00000
	1	0,66500	0,70000
	2	0,63000	0,49000
	3	0,59500	0,34300
	4	0,56000	0,24010
	5	0,52500	0,16807
	6	0,49000	0,11765
	7	0,45500	0,08235
	8	0,42000	0,05765
	9	0,38500	0,04035
	10	0,35000	0,02825
	11	0,31500	0,01977
	12	0,28000	0,01384
	13	0,24500	0,00969
	14	0,21000	0,00678
	15	0,17500	0,00475
	16	0,14000	0,00332
	17	0,10500	0,00233
	18	0,07000	0,00163
	19	0,03500	0,00114
$MaxEpoch$	20		

$$LearningRate = \alpha \times (1 - \frac{Epoch}{MaxEpoch})$$



$$LearningRate = \alpha^{Epoch}, \text{ with } 0 \leq \alpha \leq 1$$



## [recall] The overall learning process – an example

the LVQ algorithm learns the “codebook vectors” from the training data

- **[1]** choose the number of codebook vectors (cBv) to use
  - possibly (but not necessarily) an equal number of cBv for each class
- **[2]** start the learning process with a pool of cBv
  - either randomly selected from training data,
  - or randomly generated with the same scale as the training data
  - ... each cBv has the same number of attributes as data and an output class value
- **[3]** the instances in the training data are processed one-at-a-time
  - for a given training instance, select its **most similar** cBv
  - ... the selected cBv is the “winner”, also called the “best matching unit” (**BMU**)
- **[4]** the BMU gets its weights updated
  - if the BMU has the same class as the training instance, apply the **attraction** rule
  - otherwise, apply the **repulsion** rule (in relation to that training instance)

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

Codebook Vectors		
X1	X2	Y
3,58229404	0,79163723	0
7,79278348	2,33127338	0
7,93982082	2,86699026	1
3,39353321	4,67917911	1

## ... example – the first 4 dataset examples

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

Codebook Vectors		
X1	X2	Y
3,58229404	0,79163723	0
7,79278348	2,33127338	0
7,93982082	2,86699026	1
3,39353321	4,67917911	1

- ...
- [3]** the instances in the training data are processed one-at-a-time
  - for a given training instance, select its **most similar cBv**
  - ... the selected cBv is the “winner”, also called the “best matching unit” (**BMU**)
- [4]** the BMU gets its weights updated
  - if the BMU has the same class as the training instance, apply the **attraction** rule
  - otherwise, apply the **repulsion** rule (in relation to that training instance)

$$EuclideanDistance(a, b) = \sqrt{\sum_{i=1}^n (a_i - b_i)^2}$$

(fixed) LearningRate = 0,7

### Training

Codebook vectors				Input			Distances				Codebook vectors t+1			
#	X1	X2	Y	X1	X2	Y	(X1-X1)^2	(x2-X2)^2	Sum	Distance	BMU?	X1	X2	Y
1	3,58229404	0,79163723	0	3,39353321	2,33127338	0	0,035631	2,37048	2,40611	1,55118	BMU	3,450161	1,869383	0
2	7,79278348	2,33127338	0	3,39353321	2,33127338	0	19,3534	0	19,3534	4,39925		7,792783	2,331273	0
3	7,93982082	2,86699026	1	3,39353321	2,33127338	0	20,66873	0,28699	20,9557	4,57774		7,939821	2,86699	1
4	3,39353321	4,67917911	1	3,39353321	2,33127338	0	0	5,51266	5,51266	2,34791		3,393533	4,679179	1
1	3,45016146	1,86938254	0	3,11007348	1,78153964	0	0,11566	0,00772	0,12338	0,35125	BMU	3,2121	1,807893	0
2	7,79278348	2,33127338	0	3,11007348	1,78153964	0	21,92777	0,30221	22,23	4,71487		7,792783	2,331273	0
3	7,93982082	2,86699026	1	3,11007348	1,78153964	0	23,32646	1,1782	24,5047	4,95022		7,939821	2,86699	1
4	3,39353321	4,67917911	1	3,11007348	1,78153964	0	0,080349	8,39631	8,47666	2,91147		3,393533	4,679179	1
1	3,21209988	1,80789251	0	1,34380883	3,36836095	0	3,490511	2,43506	5,92557	2,43425		3,2121	1,807893	0
2	7,79278348	2,33127338	0	1,34380883	3,36836095	0	41,58927	1,07555	42,6648	6,53183		7,792783	2,331273	0
3	7,93982082	2,86699026	1	1,34380883	3,36836095	0	43,50737	0,25137	43,7587	6,61504		7,939821	2,86699	1
4	3,39353321	4,67917911	1	1,34380883	3,36836095	0	4,20137	1,71824	5,91961	2,43303	BMU	4,82834	5,596752	1
1	3,21209988	1,80789251	0	3,58229404	4,67917911	0	0,137044	8,24429	8,38133	2,89505		3,2121	1,807893	0
2	7,79278348	2,33127338	0	3,58229404	4,67917911	0	17,72822	5,51266	23,2409	4,82088		7,792783	2,331273	0
3	7,93982082	2,86699026	1	3,58229404	4,67917911	0	18,98804	3,28403	22,2721	4,71933		7,939821	2,86699	1
4	4,82834028	5,59675182	1	3,58229404	4,67917911	0	1,552631	0,84194	2,39457	1,54744	BMU	5,700573	6,239053	1

... example – for each BMU attraction or repulsion rule?

#### Codebook Vectors

X1	X2	Y
3,58229404	0,79163723	0
7,79278348	2,33127338	0
7,93982082	2,86699026	1
3,39353321	4,67917911	1

• ...

- **[4]** the BMU gets its weights updated
  - if the BMU has the same class as the training instance, apply the **attraction** rule
  - otherwise, apply the **repulsion** rule (in relation to that training instance)

#### Training

Codebook vectors			Input			Distances				Codebook vectors t+1				
#	X1	X2	Y	X1	X2	Y	$(X1-X1)^2$	$(x2-X2)^2$	Sum	Distance	BMU?	X1	X2	Y
1	3,58229404	0,79163723	0	3,39353321	2,33127338	0	0,035631	2,37048	2,40611	1,55116	BMU	3,450161	1,869383	0
2	7,79278348	2,33127338	0	3,39353321	2,33127338	0	19,3534	0	19,3534	4,39925		7,792783	2,331273	0
3	7,93982082	2,86699026	1	3,39353321	2,33127338	0	20,66873	0,28699	20,9557	4,57774		7,939821	2,86699	1
4	3,39353321	4,67917911	1	3,39353321	2,33127338	0	0	5,51266	5,51266	2,34791		3,393533	4,679179	1
1	3,45016146	1,86938254	0	3,11007348	1,78153964	0	0,11566	0,00772	0,12338	0,35125	BMU	3,2121	1,807893	0
2	7,79278348	2,33127338	0	3,11007348	1,78153964	0	21,92777	0,30221	22,23	4,71487		7,792783	2,331273	0
3	7,93982082	2,86699026	1	3,11007348	1,78153964	0	23,32646	1,1782	24,5047	4,95022		7,939821	2,86699	1
4	3,39353321	4,67917911	1	3,11007348	1,78153964	0	0,080349	8,39631	8,47666	2,91147		3,393533	4,679179	1
1	3,21209988	1,80789251	0	1,34380883	3,36836095	0	3,490511	2,43506	5,92557	2,43425		3,2121	1,807893	0
2	7,79278348	2,33127338	0	1,34380883	3,36836095	0	41,58927	1,07555	42,6648	6,53183		7,792783	2,331273	0
3	7,93982082	2,86699026	1	1,34380883	3,36836095	0	43,50737	0,25137	43,7587	6,61504		7,939821	2,86699	1
4	3,39353321	4,67917911	1	1,34380883	3,36836095	0	4,20137	1,71824	5,91961	2,43303	BMU	4,82834	5,596752	1
1	3,21209988	1,80789251	0	3,58229404	4,67917911	0	0,137044	8,24429	8,38133	2,89505		3,2121	1,807893	0
2	7,79278348	2,33127338	0	3,58229404	4,67917911	0	17,72822	5,51266	23,2409	4,82088		7,792783	2,331273	0
3	7,93982082	2,86699026	1	3,58229404	4,67917911	0	18,98804	3,28403	22,2721	4,71933		7,939821	2,86699	1
4	4,82834028	5,59675182	1	3,58229404	4,67917911	0	1,552631	0,84194	2,39457	1,54744	BMU	5,700573	6,239053	1
...														


## ... example – initially and after 1 epoch

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

(fixed) LearningRate = 0,7


Codebook Vectors		
X1	X2	Y
3,58229404	0,79163723	0
7,79278348	2,33127338	0
7,93982082	2,86699026	1
3,39353321	4,67917911	1

initially



after 1 epoch

Trained Codebook Vectors - after 1 epoch		
X1	X2	Y
2,55988367	2,54926094	0
6,04838903	3,19502377	0
7,34346105	3,5122898	1
5,70057264	6,23905272	1





## same example – but, with 100 and 500 *MaxEpoch*

time-dependent *LearningRate*

$$\text{LearningRate} = \alpha \times \left(1 - \frac{\text{Epoch}}{\text{MaxEpoch}}\right)$$

$\alpha = 0,7$ ; *MaxEpoch* = 100



**Trained Codebook Vectors - after 100 epochs**

X1	X2	Y
2,72574276	3,04268663	0
5,2636527	0,55812955	0
8,27643991	2,20561477	1
6,55207296	4,09297275	1

$\alpha = 0,7$ ; *MaxEpoch* = 500

**Trained Codebook Vectors - after 500 epochs**

X1	X2	Y
2,73469779	3,02211129	0
5,24320591	0,49988667	0
8,29041704	2,22638249	1
6,57012799	4,10547856	1



... a question!


Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

*given the dataset, the “codebook vectors” and the Euclidean distance how is the point marked (with a red diamond) classified (as a zero or as a one)?*

$\alpha = 0,7$ ;  $MaxEpoch = 500$

**Trained Codebook Vectors - after 500 epochs**

X1	X2	Y
2,73469779	3,02211129	0
5,24320591	0,49988667	0
8,29041704	2,22638249	1
6,57012799	4,10547856	1



notice that the point marked (with a red diamond) belongs to the dataset and is the instance (with class one) in the “lowest-and-most-right” location...


... a question – lets answer the question!

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

$\alpha = 0,7$ ;  $MaxEpoch = 500$

Trained Codebook Vectors - after 500 epochs

X1	X2	Y
2,73469779	3,02211129	0
5,24320591	0,49988667	0
8,29041704	2,22638249	1
6,57012799	4,10547856	1



we can identify, in the dataset, the point to be classified  
we can (visually) see which cBv are closer to the point to be classified  
so, we **just need** to calculate the **distance** (Euclidean) from **each cBv to the point...**


... the classification of the marked point!

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

$\alpha = 0,7$ ;  $MaxEpoch = 500$

Trained Codebook Vectors - after 500 epochs

X1	X2	Y
2,73469779	3,02211129	0
5,24320591	0,49988667	0
8,29041704	2,22638249	1
6,57012799	4,10547856	1



	X1	X2	class
the point to be classified	7,93982082	0,79163723	1

		X1	X2		Euclidean distance
cBv	5,24320591	0,49988667	0	2,71235144	
cBv	8,29041704	2,22638249	1	1,47696028	







so, the marked point is classified as 1 (one) why?

## an exercise – the LVQ Neural Network (NN) structure

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1

use the dataset and  
codebook to “*fill-the-gaps*”  
of the presented LVQ NN...

input instance

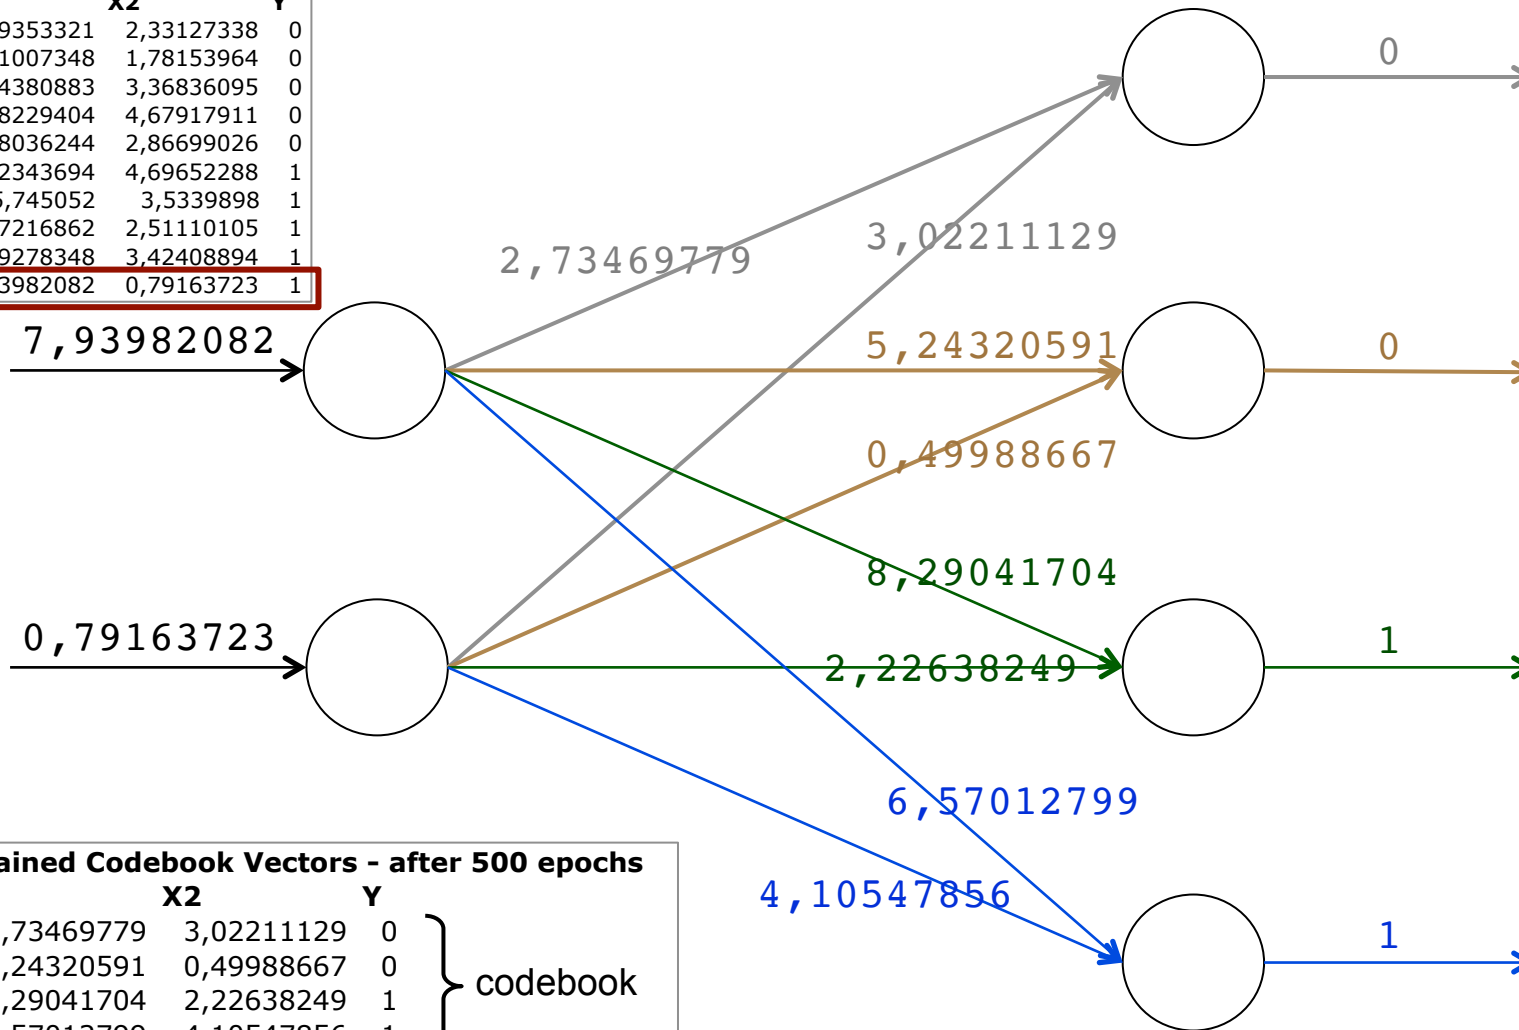


Trained Codebook Vectors - after 500 epochs

X1	X2	Y	} codebook
2,73469779	3,02211129	0	
5,24320591	0,49988667	0	
8,29041704	2,22638249	1	
6,57012799	4,10547856	1	

## ... the LVQ Neural Network NN structure

Dataset		
X1	X2	Y
3,39353321	2,33127338	0
3,11007348	1,78153964	0
1,34380883	3,36836095	0
3,58229404	4,67917911	0
2,28036244	2,86699026	0
7,42343694	4,69652288	1
5,745052	3,5339898	1
9,17216862	2,51110105	1
7,79278348	3,42408894	1
7,93982082	0,79163723	1



**Trained Codebook Vectors - after 500 epochs**

X1	X2	Y	} codebook
2,73469779	3,02211129	0	
5,24320591	0,49988667	0	
8,29041704	2,22638249	1	
6,57012799	4,10547856	1	

## LVQ – some additional approaches

### ***an alternative update-rule***

**Idea:** update not only the BMU (closest cBv), but **update the two closest cBv**

additionally we may also impose that such update only occurs  
in case the **two closest cBv represent different classes**

### ***other ideas***

### **Frequency Sensitive Competitive Learning**

the distance to a cBv is modified according to the number of data points that are assigned to this cBv

### **Size and Shape Parameters**

associate each cVb with a cluster radius; update  
radius depending on how close the data points are

### **Fuzzy LVQ**

exploits the close relationship to fuzzy clustering (an online version of it)

## Extend LVQ to deal with with categorical values...

usually, a feature may be either numeric or categorical (nominal or ordinal)

nominal – values belong to a limited and set of categories without natural ordering

e.g., “arthritis”, “asthma”, “diabetes”, “ulcers”

ordinal – values have particular order but unknown distance

e.g., “very-low”, “low”, “normal”, “high”, “very-high”

LVQ is originally designed for metric vector spaces (numeric features)

when extending LVQ to non-vector representation (i.e., to include categorical domain features) we find two main difficulties:

- a. define the distance measurement, and
- b. define the learning update rules

*extend distance measurement* – e.g., mismatch measurement on categorical features

*extend update rules* – is more complex and a “nice” approach is proposed in,  
“*Extending Learning Vector Quantization for Classifying Data with Categorical Values*”;  
by Ning Chen and Nuno Marques; *Communications in Computer and Information Science*